

Wenjie Qu

Homepage: quwenjie.github.io Email: wen_jie_qu@outlook.com

EDUCATION

National University of Singapore

Ph.D. student in Computer Science

Huazhong University of Science and Technology

B.E. in Electrical and Computer Engineering

2023.8-Now

Advisor: Prof. Jiaheng Zhang

2019.9-2023.6

RESEARCH INTERESTS

Applied Cryptography, LLM Security & Privacy.

PUBLICATION HIGHLIGHTS

7 first author papers in Top-Security/Cryptography Conferences.

13 first/co-first author papers in CCF-A Conferences.

Total Citations: **1101**, Google Scholar Link

- [1] **W. Qu***, Y. Guo*, Z. Xuan, X. Liu, J. Zhang. “Celer: A Lookup Argument for Large-Scale Queries” in *Annual International Cryptology Conference (CRYPTO)*, 2026.
- [2] **W. Qu***, Y. Guo*, Y. Ying, J. Zhang. “VerfCNN, Optimal Complexity zkSNARK for Convolutional Neural Networks” in *IEEE Symposium on Security and Privacy (SP)*, 2026
- [3] **W. Qu***, Y. Jiang*, Y. Guo, J. Zhang. “Melior, Verifiable Mixed-Integer Programming based on zkSNARKs” in *ACM Conference on Computer and Communications Security (CCS)*, 2026
- [4] Z. Qiu*, **W. Qu***, J. Zhang, B. Yuan. “V3DB: Audit-on-Demand Zero-Knowledge Proofs for Verifiable Vector Search over Committed Snapshots” in *International Conference on Very Large Data Bases (VLDB)*, 2026
- [5] X. Yang*, **W. Qu***, Y. Guo, J. Zhang. “Rarus: A Succinct and Efficient Range Proof for Polynomial-based Vector Commitment” in *USENIX Security (SEC)*, 2026
- [6] Y. Li*, **W. Qu***, L. Wu, Y. Chen, T. Cao, B. Hooi, J. Zhang. “AliMark: Enhancing Robustness of Sentence-Level Watermarking Against Text Paraphrasing” in *International Conference on Machine Learning (ICML)*, 2026
- [7] Y. Wang*, **W. Qu***, S. Zhai, Y. Jiang, Z. Liu, Y. Liu, Y. Dong, J. Zhang. “Silent Leaks: Implicit Knowledge Extraction Attack on RAG Systems” in *International Conference on Learning Representations (ICLR)*, 2026
- [8] **W. Qu**, Y. Sun, X. Liu, T. Lu, Y. Guo, K. Chen, J. Zhang. “zkGPT: An Efficient Non-interactive Zero-knowledge Proof Framework for LLM Inference” in *USENIX Security (SEC)*, 2025
- [9] **W. Qu**, W. Zheng, T. Tao, D. Yin, Y. Jiang, Z. Tian, W. Zou, J. Jia, J. Zhang. “Provably Robust Multi-bit Watermarking for AI-generated Text ” in *USENIX Security (SEC)*, 2025
- [10] **W. Qu**, Y. Zhou, Y. Wu, T. Xiao, B. Yuan, Y. Li, J. Zhang. “Prompt Inversion Attack against Collaborative Inference of Large Language Models” in *IEEE Symposium on Security and Privacy (SP)*, 2025
- [11] **W. Qu**, J. Jia, N. Gong. “REaaS: Enabling Adversarially Robust Downstream Classifiers via Robust Encoder as a Service” in *Network and Distributed System Security (NDSS)*, 2023

- [12] J. Jia*, **W. Qu***, and N. Gong. “MultiGuard: Provably Robust Multi-label Classification against Adversarial Examples” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2022, **Spotlight**
- [13] H. Wang*, **W. Qu***, G. Katz, W. Zhu, Z. Gao, H. Qiu, J. Zhuge, and C. Zhang. “jTrans: Jump-Aware Transformer for Binary Code Similarity Detection” in *International Symposium on Software Testing and Analysis (ISSTA)*, 2022

OTHER PUBLICATIONS

- [1] T. Lu, J. Zhang, Y. Guo, X. Liu, **W. Qu**, Z. Wang, W. Chen, J. Zhang. “TensorZKP: Repurposing GPU Tensor Cores for High-Performance Zero-Knowledge Proofs” in *USENIX Security (SEC)*, 2026
- [2] Y. Guo, A. Xiong, **W. Qu**, J. Zhang. “UltraProofs: Scalable Reed-Solomon Code Commitment” in *IEEE Symposium on Security and Privacy (SP)*, 2026
- [3] S. Zhao, C. Wang, Y. Li, Y. Huang, **W. Qu**, S. Lam, Y. Xie, K. Chen, J. Zhang, T. Zhang. “Towards Effective Prompt Stealing Attack against Text-to-Image Diffusion Models” in *Network and Distributed System Security (NDSS)*, 2026
- [4] J. Lu, Y. Li, T. Song, W. Wang, **W. Qu***, H. Guan, J. Zhang. “Echoes within the Reasoning: Stealthy and Effective Watermarking via Chain of Thought” in *International Conference on Machine Learning (ICML)*, 2026
- [5] Y. Guo, X. Liu, K. Huang, **W. Qu**, W. Zheng, T. Tao, D. Yin, Y. Jiang, T. Tao, J. Zhang. “DeepFold: Efficient Multilinear Polynomial Commitment from Reed-Solomon Code and Its Application to Zero-knowledge Proofs” in *USENIX Security (SEC)*, 2025
- [6] C. Li, P. Zhu, Y. Li, C. Hong, **W. Qu**, J. Zhang. “HyperPianist: Pianist with Linear-Time Prover and Logarithmic Communication Cost” in *IEEE Symposium on Security and Privacy (SP)*, 2025
- [7] K. Cheng, Y. Xia, A. Song, J. Fu, **W. Qu**, Y. Shen, J. Zhang. “Mosformer: Maliciously Secure Three-Party Inference Framework for Large Transformers” in *ACM Conference on Computer and Communications Security (CCS)*, 2025
- [8] X. Yang*, J. Han*, R Bommasani*, J Luo*, **W Qu***, W Zhou*, et al. “Reliable and Responsible Foundation Models” in *Transactions on Machine Learning Research (TMLR)*, 2025
- [9] S. Zhai, J. Li, Y. Liu, H. Chen, Z. Tian, **W. Qu**, Q. Shen, R. Jia, Y. Dong, J. Zhang. “Efficient Input-level Backdoor Defense on Text-to-Image Synthesis via Neuron Activation Variation” in *International Conference on Computer Vision (ICCV)*, 2025, **Highlight**
- [10] Z. Tian, Y. Ding, **W. Qu**, X. Yu, E. Gong, J. Zhang, J. Liu, K. Ren. “Towards Collaborative Anti-Money Laundering Among Financial Institutions” in *ACM Web Conference (WWW)*, 2025
- [11] J. Wang, **W. Qu**, Y. Rong, H. Qiu, Q. Li, Z. Li, C. Zhang. “MPass: Bypassing Learning-based Static Malware Detectors” in *Design Automation Conference (DAC)*, 2023
- [12] H. Liu*, J. Jia*, **W. Qu**, and N. Gong. “EncoderMI: Membership Inference against Pre-trained Encoders in Contrastive Learning” in *ACM Conference on Computer and Communications Security (CCS)*, 2021

SELECTED PREPRINTS

- [1] Y. Guo, **W. Qu**, L. Wu, S. Zhai, LZ. Wang, M. Xu, Y. Liu, B. Yuan, D. Song, J. Zhang. “IMMACULATE: A Practical LLM Auditing Framework via Verifiable Computation”.
- [2] **W. Qu**, X. Zhao, J. Zhang, D. Song. “Self-Sovereign Agent”.
- [3] **W. Qu**, Y. Guo, J. Zhang. “Lightning, Field-Agnostic Super-Efficient Polynomial Commitment Scheme”.
- [4] **W. Qu**, Y. Zhou, B. Wang, Y. Li, LZ. Wang, J. Jia, J. Zhang. “RepoMark: A Data-Usage Auditing Framework for Code Large Language Models”.

ACADEMIC SERVICE

Program Committee

- CCS 2025, 2026

Workshop Organizer

- ICLR 2025 Workshop on Foundation Models in the Wild

Journal Reviewer

- IoT-J, TIFS, TDSC, TMC, JSS.

Reviewer

- PKC 2024, ICML 2022, 2026, MM 2024, NeurIPS 2025, COLM 2025, 2026, CVPR 2026

SELECTED AWARDS & HONORS

- | | |
|---|------|
| • ML and Systems Rising Star (one of only two from Asian institutions; 39 recipients) | 2026 |
| • Apple AI/ML PhD Fellowship Nomination (top 3 in whole department) | 2025 |
| • Dean’s Graduate Research Excellence Award (top 5 in whole department) | 2025 |
| • DAAD AInet Fellowship | 2025 |
| • Ant Group Intech Scholarship Future (10 recipients globally) | 2025 |
| • NUS President’s Graduate Fellowship (most prestigious scholarship for NUS PhD) | 2023 |
| • Stars of Tomorrow, Microsoft Research Asia | 2023 |
| • Outstanding Graduate, HUST | 2023 |
| • Huawei Scholarship (sole recipient in department) | 2022 |
| • National Scholarship (highest honor for undergraduate in China) | 2020 |
| • Outstanding Undergraduate of Academic Performance (top 1% in department) | 2020 |